arXiv:2611.09964v1 [cs.CV] 28 Oct 2026

How Many Demonstrations Does a VLA Need? On the Sample Efficiency of VLA Adaptation

Rahul Chand
Stanford University
rahulsc@cs.stanford.edu

Miguel Villa Floran
California Polytechnic State University
miguelvf@calpoly.edu

Abstract: Pretrained vision language action (VLA) models show strong zero-shot performance, but are rarely deployed directly and are often first post-trained on demonstrations from the target setting. Yet VLAs are typically compared by zero-shot or full-data fine-tuned success, revealing little about the data and compute required for adaptation. We study adaptation efficiency for six VLAs (π_0, $\pi_{0.5}$, and GR00T N1-N1.7) on four LIBERO suites, varying both demonstrations and training steps. This reveals differences hidden by full-data fine-tuning: models within two points on the full LIBERO-Spatial dataset differ by up to 3$\times$ in success with 50 demonstrations. $\pi_{0.5}$ needs 2-5$\times$ fewer demonstrations than π_0 and is the most sample-efficient model at every budget, with the largest gap on the long-horizon suite. Within GR00T, N1.5 adapts more efficiently than N1.6 and N1.7 on three of four suites, despite their stronger zero-shot performance. We study one pretraining factor behind these differences by freezing the VLM during finetuning: this severely degrades the π models, whose VLM is jointly trained with the action expert during pretraining, but barely affects GR00T models. Finally, validation loss on held-out demonstrations tracks closed-loop success within a suite, but understates how much more sample-efficient $\pi_{0.5}$ is than π_0.

Keywords: Vision language action models, Sample efficiency, Fine-tuning, Pretraining for adaptation

1 Introduction

State-of-the-art robot manipulation policies are typically obtained in two stages. A generalist policy is first pretrained on large and diverse robot and web data, and is then fine-tuned on demonstrations collected in the target deployment setting. What ultimately determines the usefulness of a pretrained policy is therefore not only how well it performs out of the box, but how efficiently, in demonstrations and in compute, it can be adapted to a new setting. Despite this, VLAs are usually compared through zero-shot evaluations or through success after fine-tuning on a full benchmark dataset [1, 2, 3, 4]. Neither number measures the cost of adaptation. Yet many practical deployments have only a limited number of demonstrations, which makes adaptability under a demonstration budget an important measure of a pretrained policy.

This gap matters in practice. Fine-tuning on the full LIBERO-Spatial dataset compresses five of the six models we study into a 97-99% band but with 50 demonstrations, the same five models span 31-94%. Models with stronger reported zero-shot performance, trained on more data with more parameters, can match their predecessors at full data and still require up to twice the data or compute to get there.

We present a systematic study of VLA adaptation efficiency under behavior cloning, which is usually the standard way pretrained VLAs are fine-tuned on new tasks. We compare two families. π_0 [1] and $\pi_{0.5}$ [2] share the same PaliGemma VLM backbone and a flow-matching action expert, and differ primarily in their pretraining recipe. The GR00T N1, N1.5, N1.6, and N1.7 models [3, 4, 5, 6] are successive iterations, consisting of a VLM paired with a diffusion-transformer action head. Each reports

to improve on its predecessor's generalist performance. Within each family, the pretrained models act as a natural experiment over pretraining recipes, which would otherwise require prohibitively expensive pretraining ablations.

Our contributions are:

- **A large-scale study of VLA sample efficiency.** We fine-tune six VLAs on four LIBERO suites of increasing difficulty and measure adaptation along two axes: the number of demonstrations needed to reach a given success rate, and the number of optimization steps needed with the full dataset. The study comprises 120 fine-tuning runs and nearly 1400 checkpoint evaluations, more than 550k simulated episodes in total (Sections 3-4).
- **Evidence that zero-shot and full-data comparisons do not capture adaptation efficiency.** $\pi_{0.5}$ is 2–5$\times$ more demonstration-efficient than π_0 and outperforms every GR00T model at every budget, with the gap growing on the long-horizon suite. Within GR00T, N1.5 outperforms the stronger zero-shot models N1.6 and N1.7 at every tested budget on Spatial, Object, and Goal.
- **An analysis of where adaptability resides.** We study one pretraining axis, whether the VLM is trained jointly with the action head by freezing the VLM during fine-tuning. This is catastrophic for the π family but nearly free for GR00T models. (Section 5.1).
- **An evaluation of an offline proxy for adaptability.** Closed-loop budget sweeps are expensive. We test whether validation loss, which is available during training, can stand in for closed-loop success when measuring sample efficiency (Section 5.2).

2 Related Work

Pretrained VLAs. Generalist VLAs couple a pretrained VLM with an action decoder trained on large robot datasets. π_0 [1] attaches a flow-matching action expert to PaliGemma. $\pi_{0.5}$ [2] keeps this architecture and changes the pretraining recipe, co-training on heterogeneous robot, web, and high-level subtask data. GR00T N1 [3] pairs a VLM with a diffusion-transformer action head, and its successors [4, 5, 6] revise the backbone, the action head, and the training data. Other pretraining choices include action tokenization [7] and implicit world modeling [8].

Adapting VLAs. A growing body of work improves how VLAs are adapted, through better fine-tuning recipes [9], parameter-efficient tuning [10], object-centric or future-conditioned adaptation for few-shot settings [11, 12], one-shot skill adaptation [13], and regularizers that preserve pretrained representations during fine-tuning [14, 15]. These methods typically take the pretrained model as given. We study a complementary question: under behavior cloning which is usually the default adaptation procedure how strongly does the choice of pretrained model determine the data and compute needed to learn on new tasks?

Data efficiency and pretraining for adaptation. Data quantity and quality are known to dominate imitation learning outcomes [16, 17], and the choice of VLM backbone affects VLA performance [18]. In language modeling, better pretraining metrics do not always yield a model that fine-tunes better [19]. We provide a systematic measurement of this kind for robot policies, across two model families, four task suites, and both the data and compute axes of adaptation.

3 Experimental Setup

Models. We study π_0 and $\pi_{0.5}$ from the openpi release and GR00T N1, N1.5, N1.6, and N1.7, starting each run from the publicly released base checkpoint. Our default adaptation fine-tunes all parameters. In the frozen-VLM ablation (Section 5.1), we freeze the vision encoder and language model and train only the remaining parameters, chiefly the action head: about 578M of 3.50B parameters for π_0 and 693M of 3.62B for $\pi_{0.5}$, and 1.07B, 1.42B, and 1.62B for GR00T N1.5, N1.6, and N1.7.

Benchmarks. We use the four LIBERO suites [20], ordered from easier to harder: Spatial, Object, and Goal each vary a single factor (object layout, object identity, and task goal, respectively) across ten short-horizon tasks, while Long (LIBERO-10) consists of ten long-horizon tasks that chain multiple subtasks. The training sets contain 432, 454, 428, and 379 demonstrations, respectively.

Demonstration budgets. For the sample-efficiency sweeps, we fine-tune each model on the same B demonstrations from the suite's training set, with smaller budgets nested within larger ones. The π runs use seed 42 and the GR00T runs seed 43. B counts demonstrations for the whole ten-task suite, not per task. For the π models we use $B \in \{5, 10, 15, 25, 50\}$ on Spatial, Object, and Goal, and $B \in \{25, 50, 75, 100\}$ on Long, where five demonstrations per suite are not enough to make progress. Because of compute limits, GR00T N1.5, N1.6, and N1.7 are run with $B \in \{10, 15, 25, 50\}$ on Object, $B \in \{10, 25, 50\}$ on Spatial and Goal, and $B \in \{50, 100\}$ on Long. the remaining budgets are in progress and will be included in a revised version of this paper. In the *full-data* setting we fine-tune on the complete training set and track success over training steps. Full-data fine-tuning is thus the upper end of our budget axis.

Optimization. All runs use AdamW with an effective batch size of 48, so a training step processes the same number of samples for every model, and steps are directly comparable as a measure of adaptation compute. The π models use a learning rate of 5×10^{-5}. GR00T uses a learning rate and weight decay of 10^{-5} with 5% warmup, cosine decay, and gradient clipping at 1. Full-data and budget runs share these settings. The budget sweeps train for 6-8 epochs (π) and 6-7 epochs (GR00T). Full-data runs are trained until performance approximately saturates, resulting in run lengths from 1.6k steps ($\pi_{0.5}$ on Spatial, about 1.5 epochs) to 16.9k steps (π_0 on Long, 8 epochs).

Evaluation and metrics. Each checkpoint is evaluated on 400 episodes (40 per sub-task), and we report the fraction of successful episodes. We summarize each run with three quantities:

- **Best success:** the highest success rate across the evaluated checkpoints of a run.
- **Data-efficiency multiplier:** for a reference model at budget B, let B' be the smallest tested budget at which a compared model reaches at least the reference's best success. We report B/B'. Since B' is restricted to tested budgets, the multiplier is conservative and becomes a lower bound when B' is the smallest tested budget.
- **Steps to X%:** in the full-data setting, the first training step at which success is at least X% for two consecutive evaluations, for $X \in \{70, 80, 90\}$.

4 How Sample-Efficient Are Pretrained VLAs?

4.1 π_0 versus $\pi_{0.5}$

$\pi_{0.5}$ needs 2–5$\times$ fewer demonstrations. Figure 1 and Table 1 compare the two models across budgets. On every suite and at every budget, $\pi_{0.5}$ reaches a higher best success than π_0. Comparing tested budgets directly, $\pi_{0.5}$ with 25 demonstrations outperforms π_0 with 50 on Spatial (75.0% vs. 63.0%), and $\pi_{0.5}$ with 15 outperforms π_0 with 50 on Object (61.5% vs. 44.5%) and Goal (47.5% vs. 42.8%). At the low end, $\pi_{0.5}$ with only 5 demonstrations outperforms π_0 with 25 on Object (23.8% vs. 12.0%).

The gap widens on long-horizon tasks. On LIBERO-Long, $\pi_{0.5}$ with the smallest budget we tested (25 demonstrations, 63.5%) already outperforms π_0 with 100 demonstrations (49.8%), a multiplier of at least 4$\times$. Because of compute limits we did not run $\pi_{0.5}$ below 25 demonstrations on Long, so it may be even more sample-efficient there.

Full-data fine-tuning. The full training set, the upper end of our budget axis, shows the same ordering in compute (Figure 2, Appendix Table 5). The two models reach similar final performance on Spatial (97.7% vs. 97.2%) and Object (96.8% vs. 98.5%), but $\pi_{0.5}$ reaches 80% success after 0.4k

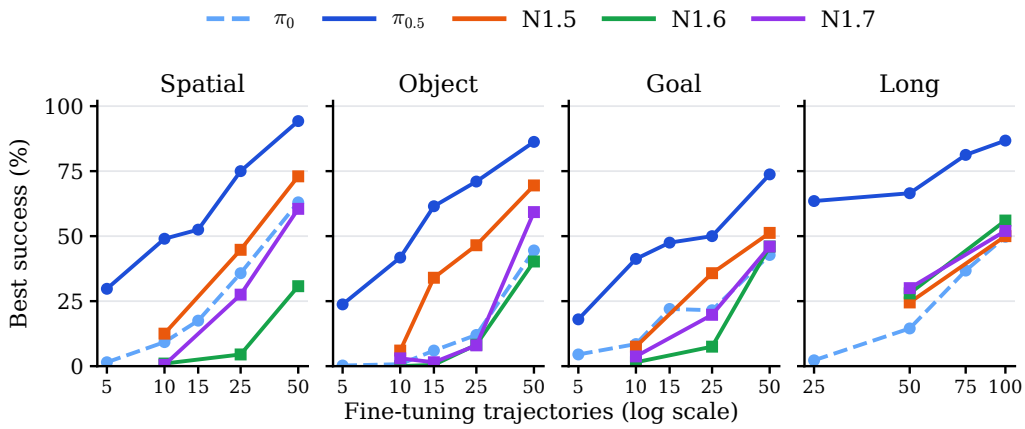

Figure 1: Best checkpoint success versus number of fine-tuning demonstrations (whole suite, log scale). A model lying above another at a smaller budget is more sample-efficient: e.g., on Spatial, $\pi_{0.5}$ with 25 demonstrations (75.0%) outperforms π_0 with 50 (63.0%). Within GR00T, N1.5 lies above N1.6 and N1.7 on Spatial, Object, and Goal.

Table 1: **How many demonstrations does $\pi_{0.5}$ need to match each model?** Each cell reads "best success of the column model with B demonstrations $\rightarrow$ smallest tested budget B' at which $\pi_{0.5}$ reaches it (B/B')." Pairwise multipliers, including within GR00T, are listed in Table 4.

Suite	B	π_0	N1.5	N1.6	N1.7
Spatial	25	35.8 $\rightarrow$ 10 (2.5×)	44.8 $\rightarrow$ 10 (2.5×)	4.5 $\rightarrow$ 5 (5.0×)	27.5 $\rightarrow$ 5 (5.0×)
Spatial	50	63.0 $\rightarrow$ 25 (2.0×)	73.0 $\rightarrow$ 25 (2.0×)	30.8 $\rightarrow$ 10 (5.0×)	60.5 $\rightarrow$ 25 (2.0×)
Object	25	12.0 $\rightarrow$ 5 (5.0×)	46.5 $\rightarrow$ 15 (1.7×)	8.2 $\rightarrow$ 5 (5.0×)	8.0 $\rightarrow$ 5 (5.0×)
Object	50	44.5 $\rightarrow$ 15 (3.3×)	69.5 $\rightarrow$ 25 (2.0×)	40.2 $\rightarrow$ 10 (5.0×)	59.2 $\rightarrow$ 15 (3.3×)
Goal	25	21.5 $\rightarrow$ 10 (2.5×)	35.8 $\rightarrow$ 10 (2.5×)	7.5 $\rightarrow$ 5 (5.0×)	19.8 $\rightarrow$ 10 (2.5×)
Goal	50	42.8 $\rightarrow$ 15 (3.3×)	51.2 $\rightarrow$ 50 (1.0×)	46.0 $\rightarrow$ 15 (3.3×)	46.0 $\rightarrow$ 15 (3.3×)
Long	50	14.5 $\rightarrow$ 25 (2.0×)	24.5 $\rightarrow$ 25 (2.0×)	28.0 $\rightarrow$ 25 (2.0×)	30.0 $\rightarrow$ 25 (2.0×)
Long	100	49.8 $\rightarrow$ 25 (4.0×)	50.0 $\rightarrow$ 25 (4.0×)	56.0 $\rightarrow$ 25 (4.0×)	52.0 $\rightarrow$ 25 (4.0×)

steps on Spatial and Object, versus 2.0k for π_0 (5×), and reaches 90% on Goal after 1.8k steps versus 6.4k (3.6×). On Long the difference extends to the ceiling: $\pi_{0.5}$ reaches 80% after 2.6k steps and peaks at 98%, whereas π_0 first reaches 80% only at its final checkpoint (16.9k steps) and peaks at 85%. A comparison of final full-data success alone would therefore miss $\pi_{0.5}$'s advantage on Spatial and Object entirely.

4.2 Stronger zero-shot models are not more sample efficient: GR00T family

Each GR00T model after N1 is reported to improve on its predecessor's generalist, zero-shot performance [4, 5, 6], with more pretraining data and progressively larger action heads (1.07B, 1.42B, and 1.62B parameters for N1.5, N1.6, and N1.7). If zero-shot capability translated into adaptability, sample efficiency should improve from N1.5 to N1.7 but it does not.

N1.5 is the most sample-efficient GR00T model on Spatial, Object, and Goal. At every tested budget on these three suites, N1.5 outperforms N1.6 and N1.7 (Figure 1. All values in Appendix Table 3). The gap is largest at small budgets: with 25 demonstrations, N1.5 reaches 44.8%, 46.5%, and 35.8% on Spatial, Object, and Goal, versus 4.5–27.5% for N1.6 and N1.7. N1.5 with 25 demonstrations even outperforms N1.6 with 50 on Spatial (44.8% vs. 30.8%) and Object (46.5% vs. 40.2%). Long is the exception, where N1.6 and N1.7 are slightly ahead at both tested budgets.

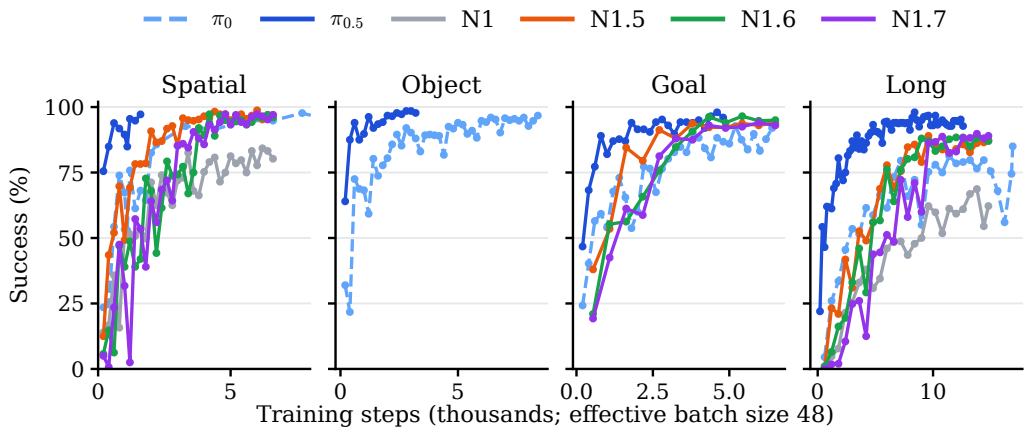

Figure 2: Full-data fine-tuning of all six models (400 evaluation episodes per checkpoint). Most models converge to similar success on the short-horizon suites, but $\pi_{0.5}$ gets there 2–5× faster than any other model. On Long it plateaus about 8 points above GR00T and 13 above π_0. GR00T full-data runs on Object are not yet available.

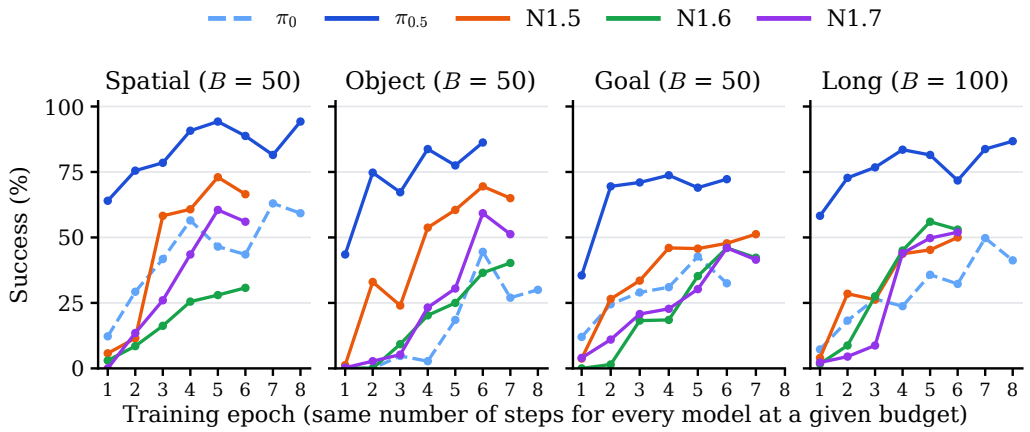

Figure 3: Success over training with a fixed demonstration budget (50. 100 for Long). At a given budget, an epoch is the same number of optimization steps for every model. $\pi_{0.5}$ leads from the first epoch on every suite. Among GR00T models, N1.5 learns fastest on Spatial, Object, and Goal. Curves for every budget are in Appendix A.

Training on full data recovers the difference. With the full training sets (Figure 2, Appendix Table 5), all three models end within a few points of each other on every suite (97–99% on Spatial, 94–97% on Goal, 88–90% on Long), but they get there at different speeds. On Spatial, N1.5 reaches 90% after 2.6k steps, whereas N1.6 and N1.7 need 4.6k and 4.2k (1.6-1.8× more). On Goal, N1.5 reaches 70% after 1.6k steps versus 2.7k for both newer models. Consistent with the budget sweeps, the newer models are competitive only on the long-horizon suite. With limited demonstrations on the short-horizon suites, the older N1.5 is clearly the better starting point.

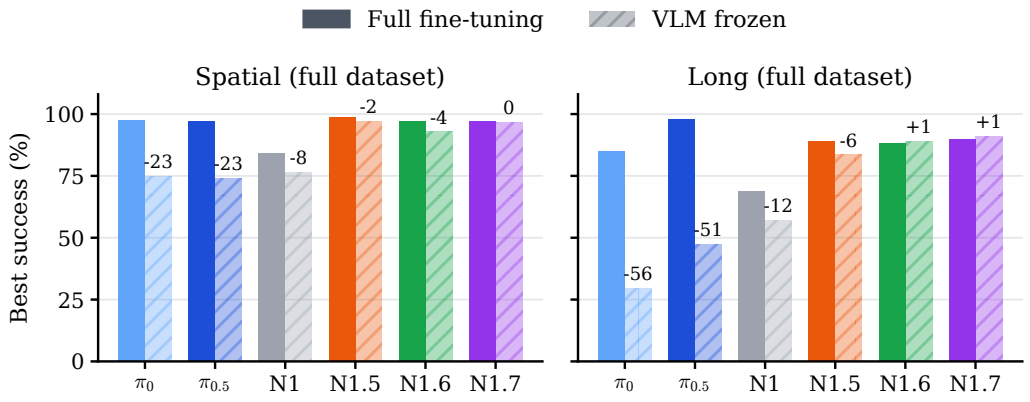

Figure 4: Best success with full fine-tuning (solid) and with the VLM frozen (hatched) on the full Spatial and Long datasets. Numbers give the change in points. Freezing the VLM costs the π models 23–56 points but changes GR00T N1.5–N1.7 by at most 6. Frozen runs train for 6 (π) and 8 (GR00T) epochs.

4.3 Across families: $\pi_{0.5}$ versus GR00T

Comparing families conflates many design choices, so we treat this comparison as indicative only. Both families share the evaluation protocol, but their learning-rate schedules differ (Section 3). $\pi_{0.5}$ outperforms every GR00T model at every tested budget on every suite (Figure 1, Table 1). The gap is largest on Long, where $\pi_{0.5}$ with 25 demonstrations outperforms every GR00T model with 100. The advantage holds in compute at a fixed budget as well (Figure 3): after two epochs, $\pi_{0.5}$ already exceeds the best success any GR00T model reaches in its whole run on every suite (e.g., 75.5% vs. 73.0% on Spatial and 72.8% vs. 56.0% on Long). With the full datasets (Figure 2, Appendix Table 5), $\pi_{0.5}$ reaches high success 2-4$\times$ faster than the fastest GR00T model on Spatial, Goal, and Long. π_0, in contrast, trails the best GR00T model at every budget. On Object, N1.5 with 25 demonstrations already outperforms π_0 with 50 (46.5% vs. 44.5%). Since π_0 shares $\pi_{0.5}$'s architecture, $\pi_{0.5}$'s lead over GR00T comes from its pretraining rather than from the π architecture.

5 Diagnosing Readiness to Adapt

Disentangling the pretraining factors that lead to these differences would require pretraining ablations beyond academic budgets. We instead ask two questions about adaptation itself: which component of the model, the VLM backbone or the action head, must be updated for adaptation to succeed, and whether adaptability can be predicted from a training signal alone, without evaluating every checkpoint in closed loop. Together, these experiments provide two practical diagnostics of adaptation: where adaptation capacity resides and whether adaptation progress can be tracked from held-out loss without closed-loop evaluation.

5.1 Freezing the VLM reveals where adaptation capacity resides

Every model we study pairs a pretrained VLM with an action head. We repeat our experiments with the VLM frozen, so that adaptation can only change the action head.

The two families place adaptation capacity in different components. Freezing the VLM is costly for the π models and nearly free for recent GR00T models (Figures 4 and 5). On Spatial, both π_0 and $\pi_{0.5}$ lose 23 points (to 74-75%). On Long, π_0 falls from 85% to 30% and $\pi_{0.5}$ from 98% to 47%. In contrast, GR00T N1.5, N1.6, and N1.7 change by between -6 and $+1$ points on both suites

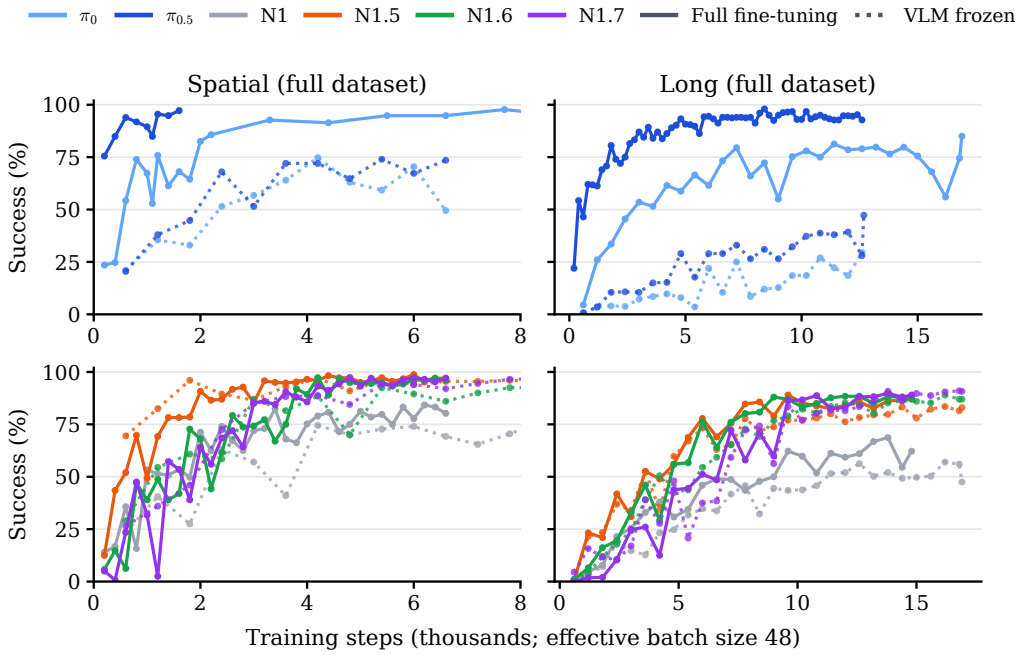

Figure 5: Full-data learning curves with full fine-tuning (solid) and the VLM frozen (dotted). Freezing slows and caps the π models, whereas the GR00T curves nearly coincide.

(the small gains are within evaluation noise), and they learn at least as quickly as with full fine-tuning (e.g., 80% on Spatial after 1.2k vs. 2.0k steps for N1.5. Appendix Table 6). GR00T N1 is in between (-8 and -12). This matches how the two families are pretrained. GR00T keeps its VLM largely frozen during pretraining [3, 4], so its action head learns to extract what it needs from fixed VLM features, and fine-tuning can do the same. The π models train the VLM jointly with the action expert [1, 2]. The two co-adapt, so downstream adaptation also relies on changing the VLM, and freezing it removes a pathway the model learned to use.

$\pi_{0.5}$'s advantage needs an adaptable backbone on Spatial, but not on Long. On Spatial, freezing the VLM removes the gap entirely (74.0% vs. 74.8% with the full dataset. Appendix Figure 9). On Long, $\pi_{0.5}$ still leads with the VLM frozen (47% vs. 30%), although both models fall well below their fully tuned performance. A deployment that can only afford to update the action head is therefore well served by GR00T but not by the π models.

5.2 Validation loss ranks budgets but understates model differences

Budget sweeps require closed-loop evaluation of every checkpoint. We therefore test whether validation loss can serve as a cheaper proxy. Specifically, we use the action MSE on held-out demonstrations, logged during each π-family budget run on Spatial, Object, and Goal. The error is computed on unnormalized actions, since normalization differs across models, and excludes the gripper, whose ± 1 values would otherwise dominate it. We compare each run's lowest validation loss with its best success (Figure 6).

Lower validation loss means higher success within a suite. Spearman ρ is -0.99 on Spatial and Object and -0.95 on Goal, and the ordering holds across models: 70 of the 75 π_0–$\pi_{0.5}$ pairs within a suite are ranked correctly. Loss scales differ between suites, however, so the pooled correlation is only -0.54.

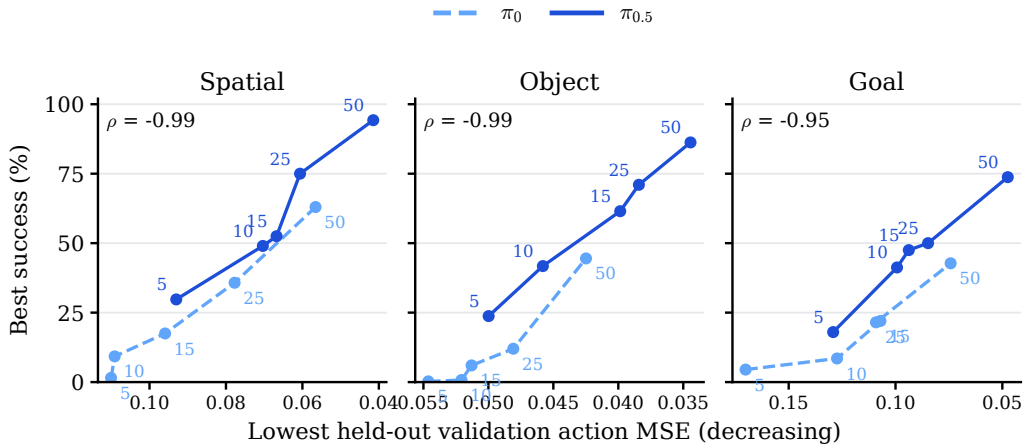

Figure 6: Lowest validation loss (held-out action MSE) versus best success for each π_0 and $\pi_{0.5}$ budget run (labels: budget). ρ: Spearman correlation over the ten runs per suite.

The size of the loss gap does not reflect the size of the efficiency gap. At equal budgets, π_0's loss is only 1.1–1.6$\times$ higher than $\pi_{0.5}$'s, regardless of whether $\pi_{0.5}$ is 2$\times$ or 5$\times$ more sample-efficient. On Object with 25 demonstrations, for example, a 1.25$\times$ loss gap corresponds to 12.0% vs. 71.0% success. Small differences in action error thus become large differences in closed-loop success. Loss is more useful as a ruler: taking the smallest $\pi_{0.5}$ budget whose loss is at least as low as π_0's recovers the data-efficiency multiplier in 8 of 12 cases, and underestimates it in the other 4 (Appendix Table 2). Validation loss is thus a useful but conservative proxy, tested here only within the π family.

6 Discussion and Conclusion

Practical implications. Freezing and held-out loss provide two inexpensive diagnostics of readiness to adapt. A frozen-VLM run tests whether adaptation can be confined to the action head: this is nearly lossless for recent GR00T models but costs the π models 23–56 success points. Within the π family, held-out action loss reliably ranks budgets within each suite ($\rho = -0.95$ to -0.99), but small loss gaps can hide large differences in closed-loop success.

Conclusion. This paper studies how efficiently pretrained VLAs adapt to new tasks under behavior cloning, across six models and four LIBERO suites. Adaptation efficiency differs several-fold between models and is predicted neither by zero-shot capability nor by full-data success. $\pi_{0.5}$ needs 2–5$\times$ fewer demonstrations than π_0, and N1.5 adapts more efficiently than its successors N1.6 and N1.7. Freezing the VLM shows that the pretraining recipe determines which components must be updated downstream, and validation loss tracks success within a suite but understates the gap between models. We therefore suggest reporting success as a function of fine-tuning demonstrations and of fine-tuning steps alongside zero-shot and full-data results.

Limitations. All experiments are conducted in simulation on LIBERO and do not include other commonly used simulation benchmarks. Real robots may behave differently, making similar experiments on physical systems an important direction for future work. Finally, released models differ in many ways at once, so our evidence on *which* pretraining choices matter is correlational.

References

[1] K. Black, N. Brown, D. Driess, A. Esmail, M. R. Equi, C. Finn, N. Fusai, L. Groom, K. Hausman, B. Ichter, S. Jakubczak, T. Jones, L. Ke, S. Levine, A. Li-Bell, M. Mothukuri, S. Nair, K. Pertsch, L. X. Shi, L. Smith, J. Tanner, Q. Vuong, A. Walling, H. Wang, and U. Zhilinsky. π_0: A Vision-Language-Action Flow Model for General Robot Control. In *Proceedings of Robotics: Science and Systems*, 2025. doi:10.15607/RSS.2025.XXI.010. URL https://www.roboticsproceedings.org/rss21/p010.html.

[2] Physical Intelligence, K. Black, N. Brown, J. Darpinian, K. Dhabalia, D. Driess, A. Esmail, M. Equi, C. Finn, N. Fusai, M. Y. Galliker, D. Ghosh, L. Groom, K. Hausman, B. Ichter, S. Jakubczak, T. Jones, L. Ke, D. LeBlanc, S. Levine, A. Li-Bell, M. Mothukuri, S. Nair, K. Pertsch, A. Z. Ren, L. X. Shi, L. Smith, J. T. Springenberg, K. Stachowicz, J. Tanner, Q. Vuong, H. Walke, A. Walling, H. Wang, L. Yu, and U. Zhilinsky. $\pi_{0.5}$: a Vision-Language-Action Model with Open-World Generalization. *arXiv preprint arXiv:2504.16054*, 2025. URL https://arxiv.org/abs/2504.16054.

[3] NVIDIA, J. Bjorck, F. Castañeda, N. Cherniadev, X. Da, R. Ding, L. J. Fan, Y. Fang, D. Fox, F. Hu, S. Huang, J. Jang, Z. Jiang, J. Kautz, K. Kundalia, L. Lao, Z. Li, Z. Lin, K. Lin, G. Liu, E. Llontop, L. Magne, A. Mandlekar, A. Narayan, S. Nasiriany, S. Reed, Y. L. Tan, G. Wang, Z. Wang, J. Wang, Q. Wang, J. Xiang, Y. Xie, Y. Xu, Z. Xu, S. Ye, Z. Yu, A. Zhang, H. Zhang, Y. Zhao, R. Zheng, and Y. Zhu. GR00T N1: An Open Foundation Model for Generalist Humanoid Robots. *arXiv preprint arXiv:2503.14734*, 2025. URL https://arxiv.org/abs/2503.14734.

[4] J. Bjorck, F. Castañeda, N. Cherniadev, X. Da, R. Ding, L. J. Fan, Y. Fang, D. Fox, F. Hu, S. Huang, J. Jang, Z. Jiang, J. Kautz, K. Kundalia, Z. Li, Z. Lin, K. Lin, G. Liu, L. Magne, A. Mandlekar, A. Narayan, S. Nasiriany, S. Reed, Y. L. Tan, G. Wang, J. Wang, Q. Wang, J. Xiang, Y. Xie, Y. Xu, S. Ye, Z. Yu, A. Zhang, H. Zhang, Y. Zhao, R. Zheng, Y. Zhu, and X. Jiang. GR00T N1.5: An Improved Open Foundation Model for Generalist Humanoid Robots. NVIDIA Research technical report, June 2025. URL https://research.nvidia.com/labs/gear/gr00t-n15/. Published June 9, 2025.

[5] NVIDIA GEAR Lab. GR00T N1.6: An Improved Open Foundation Model for Generalist Humanoid Robots. NVIDIA Research blog post, Dec. 2025. URL https://research.nvidia.com/labs/gear/gr00t-n1_6/. Published December 15, 2025.

[6] NVIDIA. Isaac GR00T N1.7. GitHub repository, 2026. URL https://github.com/NVIDIA/Isaac-GR00T.

[7] K. Pertsch, K. Stachowicz, B. Ichter, D. Driess, S. Nair, Q. Vuong, O. Mees, C. Finn, and S. Levine. FAST: Efficient Action Tokenization for Vision-Language-Action Models. In *Proceedings of Robotics: Science and Systems*, 2025. doi:10.15607/RSS.2025.XXI.012. URL https://www.roboticsproceedings.org/rss21/p012.html.

[8] R. Zheng, J. Wang, S. Reed, J. Bjorck, Y. Fang, F. Hu, J. Jang, K. Kundalia, Z. Lin, L. Magne, A. Narayan, Y. L. Tan, G. Wang, Q. Wang, J. Xiang, Y. Xu, S. Ye, J. Kautz, F. Huang, Y. Zhu, and L. Fan. FLARE: Robot Learning with Implicit World Modeling. In *Proceedings of The 9th Conference on Robot Learning*, volume 305 of *Proceedings of Machine Learning Research*, pages 3952–3971. PMLR, 2025. URL https://proceedings.mlr.press/v305/zheng25a.html.

[9] M. J. Kim, C. Finn, and P. Liang. Fine-Tuning Vision-Language-Action Models: Optimizing Speed and Success. In *Proceedings of Robotics: Science and Systems*, 2025. doi:10.15607/RSS.2025.XXI.017. URL https://www.roboticsproceedings.org/rss21/p017.html.

[10] F. Ferchau, D. Pommer, and C. Axenie. On the Efficiency of LoRA Fine-Tuning for Vision-Language-Action Models in Industrial Robotic Manipulation. *arXiv preprint arXiv:2607.10172*, 2026. URL https://arxiv.org/abs/2607.10172.

[11] P. Li, Y. Wu, Z. Xi, W. Li, Y. Huang, Z. Zhang, Y. Chen, J. Wang, S.-C. Zhu, T. Liu, and S. Huang. ControlVLA: Few-shot Object-centric Adaptation for Pre-trained Vision-Language-Action Models. In *Proceedings of The 9th Conference on Robot Learning*, volume 305 of *Proceedings of Machine Learning Research*, pages 1898–1913. PMLR, 2025. URL https://proceedings.mlr.press/v305/li25c.html.

[12] D. M. Nguyen, N. T. Diep, B. G. Nguyen, T.-B. Ho, D. Le, T. Q. Nguyen, T.-L. Ha, N. Tran, B. Thach, N. X. Tran, T. A. Tran, A. Habuda, P. L. Møller, T. N. Le, D. Sonntag, M. Niepert, K. D. Doan, V. Duong, H. Ngo, M. N. Vu, D. M. H. Nguyen, A. T. Le, and N. A. Vien. FOCA: Future-Oriented Conditioning for Data-Efficient Vision-Language-Action Adaptation. *arXiv preprint arXiv:2606.20867*, 2026. URL https://arxiv.org/abs/2606.20867.

[13] R. Zhao, T. Ingebrand, S. Chinchali, and U. Topcu. MoS-VLA: A Vision-Language-Action Model with One-Shot Skill Adaptation. *arXiv preprint arXiv:2510.16617*, 2025. URL https://arxiv.org/abs/2510.16617.

[14] C. Huang, M. M. Zhang, R. Azarcon, G. Chou, and Z. Kira. MAPS: Preserving vision-language representations via module-wise proximity scheduling for better VLA generalization. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2026. URL https://mapsvla.github.io/.

[15] N. Kachaev, M. Kolosov, D. Zelezetsky, A. K. Kovalev, and A. I. Panov. Don't Blind Your VLA: Aligning Visual Representations for OOD Generalization. *arXiv preprint arXiv:2510.25616*, 2025. URL https://arxiv.org/abs/2510.25616.

[16] A. Mandlekar, D. Xu, J. Wong, S. Nasiriany, C. Wang, R. Kulkarni, L. Fei-Fei, S. Savarese, Y. Zhu, and R. Martín-Martín. What Matters in Learning from Offline Human Demonstrations for Robot Manipulation. In *Proceedings of the 5th Conference on Robot Learning*, volume 164 of *Proceedings of Machine Learning Research*, pages 1678–1690. PMLR, 2022. URL https://proceedings.mlr.press/v164/mandlekar22a.html.

[17] F. Lin, Y. Hu, P. Sheng, C. Wen, J. You, and Y. Gao. Data Scaling Laws in Imitation Learning for Robotic Manipulation. In *International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/88b7b2c896506daabc8d3fd587055167-Abstract-Conference.html.

[18] J. Zhang, X. Chen, Q. Wang, M. Li, Y. Guo, Y. Hu, J. Zhang, S. Bai, J. Lin, and J. Chen. VLM4VLA: Revisiting Vision-Language-Models in Vision-Language-Action Models. *arXiv preprint arXiv:2601.03309*, 2026. URL https://arxiv.org/abs/2601.03309.

[19] J. M. Springer, S. Goyal, K. Wen, T. Kumar, X. Yue, S. Malladi, G. Neubig, and A. Raghunathan. Overtrained Language Models Are Harder to Fine-Tune. In *International Conference on Machine Learning*, 2025. URL https://arxiv.org/abs/2503.19206.

[20] B. Liu, Y. Zhu, C. Gao, Y. Feng, Q. Liu, Y. Zhu, and P. Stone. LIBERO: Benchmarking Knowledge Transfer for Lifelong Robot Learning. In *Advances in Neural Information Processing Systems*, volume 36, 2023. URL https://proceedings.nips.cc/paper_files/paper/2023/hash/8c3c666820ea055a77726d66fc7d447f-Abstract-Datasets_and_Benchmarks.html.

A Learning Curves

Figures 7–8 show success over training epochs for every budget, from which the best-checkpoint values in the main text are taken.

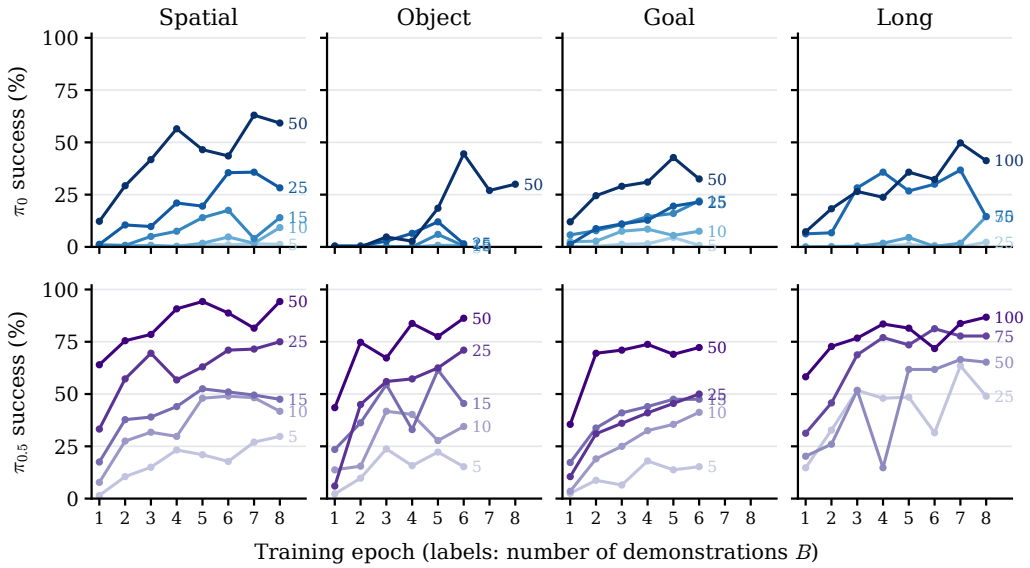

Figure 7: π_0 (top) and $\pi_{0.5}$ (bottom) budget sweeps with full fine-tuning, 400 evaluation episodes per checkpoint. Darker lines denote larger budgets.

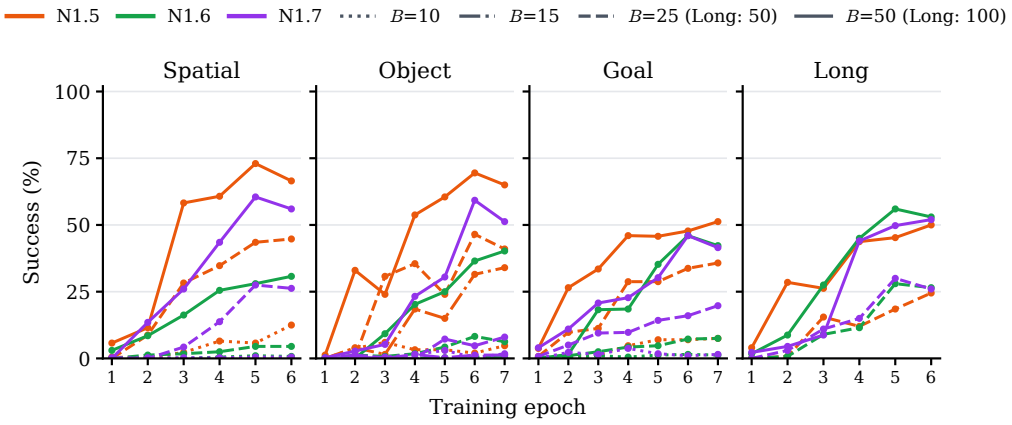

Figure 8: GR00T budget sweeps with full fine-tuning, 400 evaluation episodes per checkpoint. Colors denote models and line styles budgets.

B Fine-Tuning with the VLM Frozen

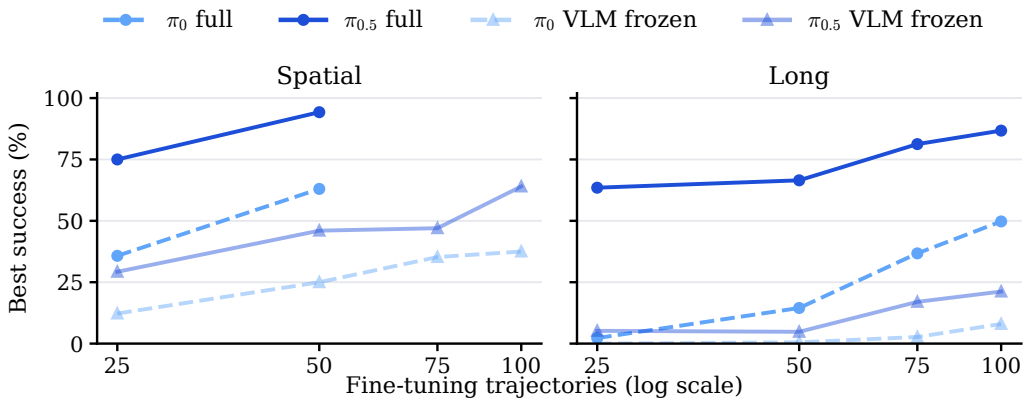

Figure 9: Best success versus fine-tuning demonstrations for π_0 and $\pi_{0.5}$ with full fine-tuning (opaque) and the VLM frozen (faded).

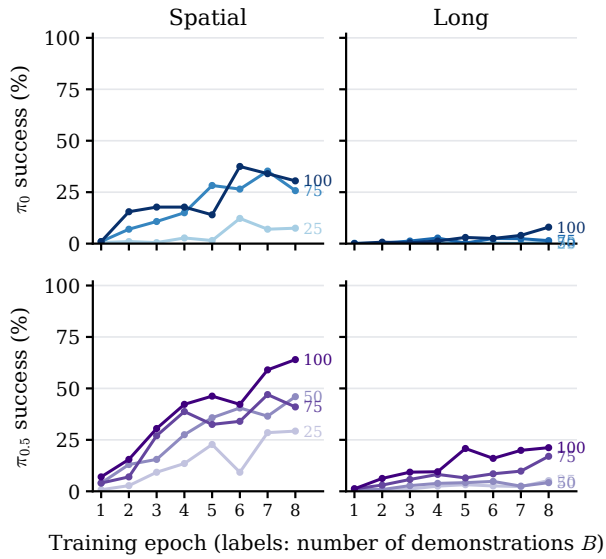

Figure 10: π_0 (top) and $\pi_{0.5}$ (bottom) budget sweeps with the VLM frozen. Only the best success (25.0%) is available for the π_0 Spatial run with 50 demonstrations, so its curve is not shown.

C Validation Loss as a Proxy

Table 2: Lowest validation loss (action MSE) and best success (%) for every π_0 and $\pi_{0.5}$ budget run. MSE ratio: π_0 over $\pi_{0.5}$ at the same budget. The multiplier from success is defined as in Section 3. The multiplier from MSE uses the smallest $\pi_{0.5}$ budget whose validation loss is at most π_0's. The MSE-based multiplier matches the success-based one in 8 of 12 cases and is never larger.

		π_0		$\pi_{0.5}$		MSE	Multiplier from	
Suite	B	MSE	Success	MSE	Success	ratio	success	MSE
Spatial	5	0.1101	1.5	0.0930	29.8	1.18	–	–
Spatial	10	0.1091	9.2	0.0703	49.0	1.55	2.0×	2.0×
Spatial	15	0.0959	17.5	0.0668	52.5	1.44	3.0×	3.0×
Spatial	25	0.0777	35.8	0.0606	75.0	1.28	2.5×	2.5×
Spatial	50	0.0566	63.0	0.0415	94.2	1.36	2.0×	1.0×
Object	5	0.0546	0.2	0.0500	23.8	1.09	–	–
Object	10	0.0521	0.8	0.0458	41.8	1.14	2.0×	2.0×
Object	15	0.0513	6.0	0.0399	61.5	1.29	3.0×	3.0×
Object	25	0.0481	12.0	0.0384	71.0	1.25	5.0×	2.5×
Object	50	0.0425	44.5	0.0344	86.2	1.23	3.3×	3.3×
Goal	5	0.1702	4.5	0.1292	18.0	1.32	–	–
Goal	10	0.1274	8.5	0.0993	41.2	1.28	2.0×	1.0×
Goal	15	0.1071	22.0	0.0938	47.5	1.14	1.5×	1.5×
Goal	25	0.1091	21.5	0.0848	50.0	1.29	2.5×	2.5×
Goal	50	0.0742	42.8	0.0475	73.8	1.56	3.3×	1.0×

D Best Success at Every Budget

Table 3: Best checkpoint success (%) at every fine-tuning budget B (demonstrations per suite). N1.5–N1.7 denote GR00T models. "–" marks budgets not run.

Spatial	B=5	B=10	B=15	B=25	B=50
π_0	1.5	9.2	17.5	35.8	63.0
$\pi_{0.5}$	29.8	49.0	52.5	75.0	94.2
N1.5	–	12.5	–	44.8	73.0
N1.6	–	1.0	–	4.5	30.8
N1.7	–	0.5	–	27.5	60.5
Object	B=5	B=10	B=15	B=25	B=50
π_0	0.2	0.8	6.0	12.0	44.5
$\pi_{0.5}$	23.8	41.8	61.5	71.0	86.2
N1.5	–	6.0	34.0	46.5	69.5
N1.6	–	0.0	0.5	8.2	40.2
N1.7	–	3.0	1.5	8.0	59.2
Goal	B=5	B=10	B=15	B=25	B=50
π_0	4.5	8.5	22.0	21.5	42.8
$\pi_{0.5}$	18.0	41.2	47.5	50.0	73.8
N1.5	–	7.5	–	35.8	51.2
N1.6	–	1.5	–	7.5	46.0
N1.7	–	3.8	–	19.8	46.0
Long	B=25	B=50	B=75	B=100	
π_0	2.2	14.5	36.8	49.8	
$\pi_{0.5}$	63.5	66.5	81.2	86.8	
N1.5	–	24.5	–	50.0	
N1.6	–	28.0	–	56.0	
N1.7	–	30.0	–	52.0	
Spatial (VLM frozen)	B=25	B=50	B=75	B=100	
π_0	12.2	–	35.2	37.5	
$\pi_{0.5}$	29.2	46.0	47.0	64.0	
Long (VLM frozen)	B=25	B=50	B=75	B=100	
π_0	0.0	0.5	2.7	8.0	
$\pi_{0.5}$	5.2	4.8	17.0	21.2	

Table 4: All pairwise data-efficiency multipliers. For a reference model at budget B, B' is the smallest tested budget at which the compared model reaches at least the reference's best success. Pairs where the compared model needs the full budget B are omitted.

	Reference model			Compared model			
Suite	Model	B	Best (%)	Model	B'	Best (%)	B/B'
Spatial	π_0	25	35.8	$\pi_{0.5}$	10	49.0	2.5×
Spatial	π_0	50	63.0	$\pi_{0.5}$	25	75.0	2.0×
Spatial	N1.5	25	44.8	$\pi_{0.5}$	10	49.0	2.5×
Spatial	N1.5	50	73.0	$\pi_{0.5}$	25	75.0	2.0×
Spatial	N1.6	25	4.5	$\pi_{0.5}$	5	29.8	5.0×
Spatial	N1.6	50	30.8	$\pi_{0.5}$	10	49.0	5.0×
Spatial	N1.7	25	27.5	$\pi_{0.5}$	5	29.8	5.0×
Spatial	N1.7	50	60.5	$\pi_{0.5}$	25	75.0	2.0×
Spatial	N1.6	25	4.5	N1.5	10	12.5	2.5×
Spatial	N1.6	50	30.8	N1.5	25	44.8	2.0×
Object	π_0	25	12.0	$\pi_{0.5}$	5	23.8	5.0×
Object	π_0	50	44.5	$\pi_{0.5}$	15	61.5	3.3×
Object	N1.5	25	46.5	$\pi_{0.5}$	15	61.5	1.7×
Object	N1.5	50	69.5	$\pi_{0.5}$	25	71.0	2.0×
Object	N1.6	25	8.2	$\pi_{0.5}$	5	23.8	5.0×
Object	N1.6	50	40.2	$\pi_{0.5}$	10	41.8	5.0×
Object	N1.7	25	8.0	$\pi_{0.5}$	5	23.8	5.0×
Object	N1.7	50	59.2	$\pi_{0.5}$	15	61.5	3.3×
Object	N1.6	25	8.2	N1.5	15	34.0	1.7×
Object	N1.6	50	40.2	N1.5	25	46.5	2.0×
Object	N1.7	25	8.0	N1.5	15	34.0	1.7×
Goal	π_0	25	21.5	$\pi_{0.5}$	10	41.2	2.5×
Goal	π_0	50	42.8	$\pi_{0.5}$	15	47.5	3.3×
Goal	N1.5	25	35.8	$\pi_{0.5}$	10	41.2	2.5×
Goal	N1.6	25	7.5	$\pi_{0.5}$	5	18.0	5.0×
Goal	N1.6	50	46.0	$\pi_{0.5}$	15	47.5	3.3×
Goal	N1.7	25	19.8	$\pi_{0.5}$	10	41.2	2.5×
Goal	N1.7	50	46.0	$\pi_{0.5}$	15	47.5	3.3×
Goal	N1.6	25	7.5	N1.5	10	7.5	2.5×
Long	π_0	50	14.5	$\pi_{0.5}$	25	63.5	2.0×
Long	π_0	75	36.8	$\pi_{0.5}$	25	63.5	3.0×
Long	π_0	100	49.8	$\pi_{0.5}$	25	63.5	4.0×
Long	N1.5	50	24.5	$\pi_{0.5}$	25	63.5	2.0×
Long	N1.5	100	50.0	$\pi_{0.5}$	25	63.5	4.0×
Long	N1.6	50	28.0	$\pi_{0.5}$	25	63.5	2.0×
Long	N1.6	100	56.0	$\pi_{0.5}$	25	63.5	4.0×
Long	N1.7	50	30.0	$\pi_{0.5}$	25	63.5	2.0×
Long	N1.7	100	52.0	$\pi_{0.5}$	25	63.5	4.0×

E Full-Data Adaptation Speed

Table 5: Full-data fine-tuning: training steps (effective batch size 48) until success is at least 70/80/90% at two consecutive evaluations, and best success. N1–N1.7 denote GR00T models. "–" means the threshold was not reached within the run. To save compute, some runs are evaluated less frequently (every 200–1,100 steps), which bounds the precision of these numbers.

Suite	Model	70%	80%	90%	Best (%)
Spatial	π_0	2.0k	2.0k	3.3k	97.7
	$\pi_{0.5}$	0.2k	0.4k	0.6k	97.2
	N1	3.0k	6.2k	–	84.2
	N1.5	1.4k	2.0k	2.6k	98.8
	N1.6	2.6k	3.8k	4.6k	97.2
	N1.7	3.0k	3.0k	4.2k	97.2
Object	π_0	1.4k	2.0k	2.8k	96.8
	$\pi_{0.5}$	0.4k	0.4k	1.2k	98.5
Goal	π_0	2.2k	3.0k	6.4k	92.8
	$\pi_{0.5}$	0.6k	0.8k	1.8k	98.0
	N1.5	1.6k	2.7k	3.8k	94.0
	N1.6	2.7k	3.3k	3.8k	96.5
	N1.7	2.7k	2.7k	4.3k	94.0
Long	π_0	6.6k	16.9k	–	85.0
	$\pi_{0.5}$	1.6k	2.6k	4.6k	98.0
	N1	–	–	–	68.7
	N1.5	7.2k	7.8k	–	89.0
	N1.6	7.2k	7.8k	–	88.4
	N1.7	9.6k	9.6k	–	89.8

Table 6: Frozen-VLM fine-tuning on the full datasets: training steps until success is at least 70/80/90% at two consecutive evaluations, and best success. Same conventions as Table 5.

Suite	Model	70%	80%	90%	Best (%)
Spatial	π_0	–	–	–	74.8
	$\pi_{0.5}$	3.6k	–	–	74.0
	N1	4.2k	–	–	76.5
	N1.5	1.2k	1.2k	3.6k	97.0
	N1.6	3.0k	3.0k	7.2k	92.8
	N1.7	2.4k	3.0k	5.4k	96.8
Long	π_0	–	–	–	29.5
	$\pi_{0.5}$	–	–	–	47.2
	N1	–	–	–	57.0
	N1.5	7.8k	15.6k	–	83.5
	N1.6	9.0k	10.2k	–	89.0
	N1.7	7.8k	10.8k	16.2k	91.0